

CARL: Calibration Aware Residual Learning

Jonathan Y. Ma

June 9, 2026

Contents

I	Foundations	3
0.1	Introduction	4
0.2	Literature Review	4
0.2.1	Forecasting Foundations	4
0.2.2	Diffusion Models and Generative Modeling	4
0.2.3	Calibration and Reliability	5
0.2.4	Uncertainty Quantification	6
0.2.5	Conformal Prediction	6
0.2.6	Proper Scoring Rules and Forecast Evaluation	6
0.2.7	Residual Learning and Error Modeling	6
0.2.8	Model Misspecification and Robustness	6
0.3	Background	7
0.3.1	Forecasting Models	7
0.3.2	Residual Processes	7
0.3.3	Diffusion Models	7
0.3.4	Calibration	7
0.3.5	Proper Scoring Rules	7
0.3.6	Uncertainty Quantification	7
II	Calibration Aware Diffusion Residual Learning	8
0.4	Problem Formulation	9
0.4.1	Forecasting System	9
0.4.2	Standard Residual Learning	9
0.4.3	Proposed Framework	9
0.5	Calibration Objectives	9
0.5.1	Coverage Loss	9
0.5.2	PIT-Based Calibration Loss	9
0.5.3	ECE-Based Loss	9
0.5.4	CRPS-Based Loss	9
0.5.5	Combined Objective	9
0.6	CADRL Training Procedure	10

<i>CONTENTS</i>	2
III Forecasting Testbed	11
0.7 Forecast Models	12
0.7.1 Classical Benchmark	12
0.7.2 Deep Benchmark	12
IV Numerical Experiments	13
0.8 Controlled Innovation Misspecification	14
V Evaluation	15
0.9 Calibration Evaluation	16
0.10 Robustness Analysis	16
VI Real Data Case Study	17
0.11 Empirical Application	18

Part I

Foundations

0.1 Introduction

Motivation

Forecast uncertainty depends heavily on residual behavior.

Diffusion models can learn complex residual distributions.

Existing diffusion methods do not optimize for calibration.

Calibration is essential for trustworthy uncertainty quantification.

A calibration-aware training paradigm may improve uncertainty quality.

Contribution

Propose Calibration-Aware Diffusion Residual Learning.

Incorporate calibration into diffusion training.

Compare calibration-aware and standard diffusion learning.

Evaluate performance under controlled innovation misspecification.

Demonstrate effectiveness across both classical and deep forecasting systems.

0.2 Literature Review

0.2.1 Forecasting Foundations

Lütkepohl (2005): New Introduction to Multiple Time Series Analysis

[Link](#)

Karlsson (2013): Forecasting with Bayesian Vector Autoregressions

[Link](#)

Salinas et al. (2020): DeepAR: Probabilistic Forecasting with Autoregressive Recurrent Networks

[Link](#)

0.2.2 Diffusion Models and Generative Modeling

Ho et al. (2020): Denoising Diffusion Probabilistic Models

[Link](#)

Song et al. (2021): Score-Based Generative Modeling through Stochastic Differential Equations

[Link](#)

Rasul et al. (2021): Autoregressive Denoising Diffusion Models for Multivariate Probabilistic Time Series Forecasting (TimeGrad)

[Link](#)

Yan et al. (2021): ScoreGrad: Multivariate Probabilistic Time Series Forecasting with Continuous Energy-Based Generative Models

[Link](#)

0.2.3 Calibration and Reliability

Dawid (1982): The Well-Calibrated Bayesian

[Link](#)

Guo et al. (2017): On Calibration of Modern Neural Networks

[Link](#)

Huang, Park, and Simeone (2023): Calibration-Aware Bayesian Learning

[Link](#)

Bohdal et al. (2021): Meta-Calibration: Learning of Model Calibration Using Differentiable Expected Calibration Error

[Link](#)

Kumar et al. (2018): Trainable Calibration Measures for Neural Networks from Kernel Mean Embeddings

[Link](#)

Detommaso et al. (2022): Uncertainty Calibration in Bayesian Neural Networks via Distance-Aware Priors

[Link](#)

0.2.4 Uncertainty Quantification

Hüllermeier and Waegeman (2021): Aleatoric and Epistemic Uncertainty in Machine Learning

[Link](#)

0.2.5 Conformal Prediction

Angelopoulos and Bates (2021): A Gentle Introduction to Conformal Prediction and Distribution-Free Uncertainty Quantification

[Link](#)

0.2.6 Proper Scoring Rules and Forecast Evaluation

Gneiting and Raftery (2007): Strictly Proper Scoring Rules, Prediction, and Estimation

[Link](#)

Dawid (1984): Present Position and Potential Developments: Some Personal Views on Probabilistic Forecasting

[Link](#)

Gneiting, Balabdaoui, and Raftery (2007): Probabilistic Forecasts, Calibration and Sharpness

[Link](#)

0.2.7 Residual Learning and Error Modeling

Friedman (2001): Greedy Function Approximation: A Gradient Boosting Machine

[Link](#)

He et al. (2016): Deep Residual Learning for Image Recognition

[Link](#)

0.2.8 Model Misspecification and Robustness

Masegosa (2019): Learning under Model Misspecification: Applications to Variational and Ensemble methods

[Link](#)

0.3 Background

0.3.1 Forecasting Models

VAR

RNN

0.3.2 Residual Processes

Residuals as empirical innovation proxies

0.3.3 Diffusion Models

DDPM

Score-based Models

SDE Interpretation

0.3.4 Calibration

Coverage

Reliability

PIT and ECE

0.3.5 Proper Scoring Rules

CRPS

Energy Score

Log Predictive Density

0.3.6 Uncertainty Quantification

Predictive distributions

Forecast trajectories

probabilistic forecasting

Part II

Calibration Aware Diffusion Residual Learning

0.4 Problem Formulation

0.4.1 Forecasting System

Given:

$$\hat{y}_t = f(y_{1:t-1})$$

compute residuals: $r_t = y_t - \hat{y}_t$ and learn: $p(r)$ using diffusion models.

0.4.2 Standard Residual Learning

Current diffusion models optimize: $\mathcal{L}_{\text{diffusion}}$ through score matching or denoising.

Calibration is evaluated only after training.

0.4.3 Proposed Framework

Calibration becomes part of the optimization objective: $\mathcal{L} = \mathcal{L}_{\text{diffusion}} + \lambda \mathcal{L}_{\text{calibration}}$

0.5 Calibration Objectives

0.5.1 Coverage Loss

Penalize deviations from nominal coverage:

$$\mathcal{L}_{\text{coverage}} = (\hat{C} - (1 - \alpha))^2$$

0.5.2 PIT-Based Calibration Loss

Encourage: $u_i = F_i(y_i)$ to approach $\text{Uniform}(0,1)$.

0.5.3 ECE-Based Loss

Penalize confidence-frequency mismatch.

0.5.4 CRPS-Based Loss

Use proper scoring rules as calibration objectives.

0.5.5 Combined Objective

Investigate: $\mathcal{L} = \mathcal{L}_{\text{diffusion}} + \lambda \mathcal{L}_{\text{calibration}}$ for different calibration choices.

0.6 CADRL Training Procedure

Step 1: Fit forecasting model.

Step 2: Extract residuals.

Step 3: Train diffusion residual model.

Step 4: Generate predictive residual samples.

Step 5: Compute predictive distributions.

Step 6: Evaluate calibration.

Step 7: Backpropagate calibration loss.

Step 8: Update diffusion model.

Part III

Forecasting Testbed

0.7 Forecast Models

0.7.1 Classical Benchmark

VAR

Motivation: interpretable, linear, statistical baseline

0.7.2 Deep Benchmark

RNN / DeepAR

Motivation: nonlinear forecasting, representation learning, modern forecasting baseline

Evaluate whether calibration-aware residual learning is effective across fundamentally different forecasting architectures.

Part IV

Numerical Experiments

0.8 Controlled Innovation Misspecification

Determine when calibration-aware diffusion learning improves uncertainty quantification.

DGP 1: Gaussian Innovations $u_t \sim N(0, \Sigma)$ Baseline.

DGP 2: Moderate Heavy Tails $u_t \sim t_5$

DGP 3: Extreme Heavy Tails $u_t \sim t_3$

DGP 4: Mixture Innovations $u_t \sim 0.8N(0, 1) + 0.2N(0, 9)$

DGP 5: Heteroskedastic Innovations $u_t \sim N(0, \sigma_t^2)$

Experimental Comparisons

Forecast Models: VAR, RNN

Residual Models: Gaussian, Student-t, Bootstrap, Standard Diffusion, Calibration-Aware Diffusion

Training Schemes

Standard $\mathcal{L} = \mathcal{L}_{\text{diffusion}}$ Proposed $\mathcal{L} = \mathcal{L}_{\text{diffusion}} + \lambda \mathcal{L}_{\text{calibration}}$

Part V

Evaluation

0.9 Calibration Evaluation

Coverage: Empirical Coverage, Coverage Error

Reliability: ECE, Reliability Diagrams, PIT Histograms

Probabilistic Forecast Quality: CRPS, Energy Score, Log Predictive Density

Sharpness: Average Interval Width, Sharpness Score

0.10 Robustness Analysis

Distributional Perturbations

Investigate: variance inflation, tail amplification, mixture contamination, regime shifts

Metrics, calibration degradation, coverage degradation, sharpness degradation

Part VI

Real Data Case Study

0.11 Empirical Application

Candidate Domain: Financial returns